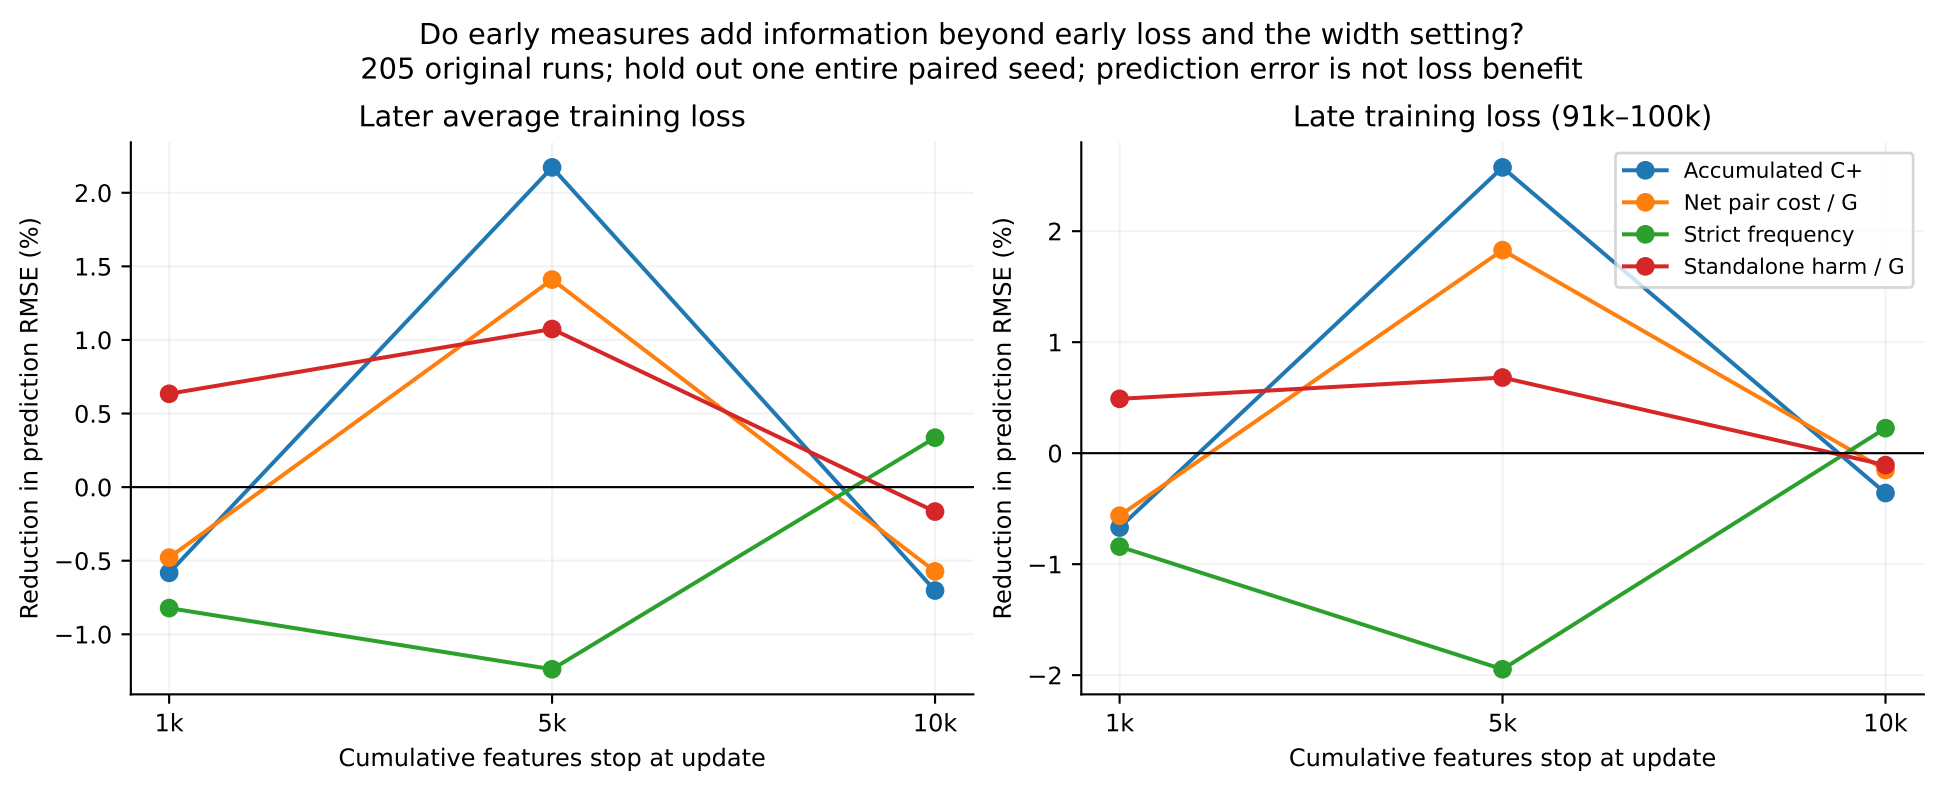
Do early measures add information beyond early loss and the width setting?
205 original runs; hold out one entire paired seed; prediction error is not loss benefit
Later average training loss
Late training loss (91k–100k)
Reduction in prediction RMSE (%)
Cumulative features stop at update
1k
5k
10k
Accumulated C+
Net pair cost / G
Strict frequency
Standalone harm / G