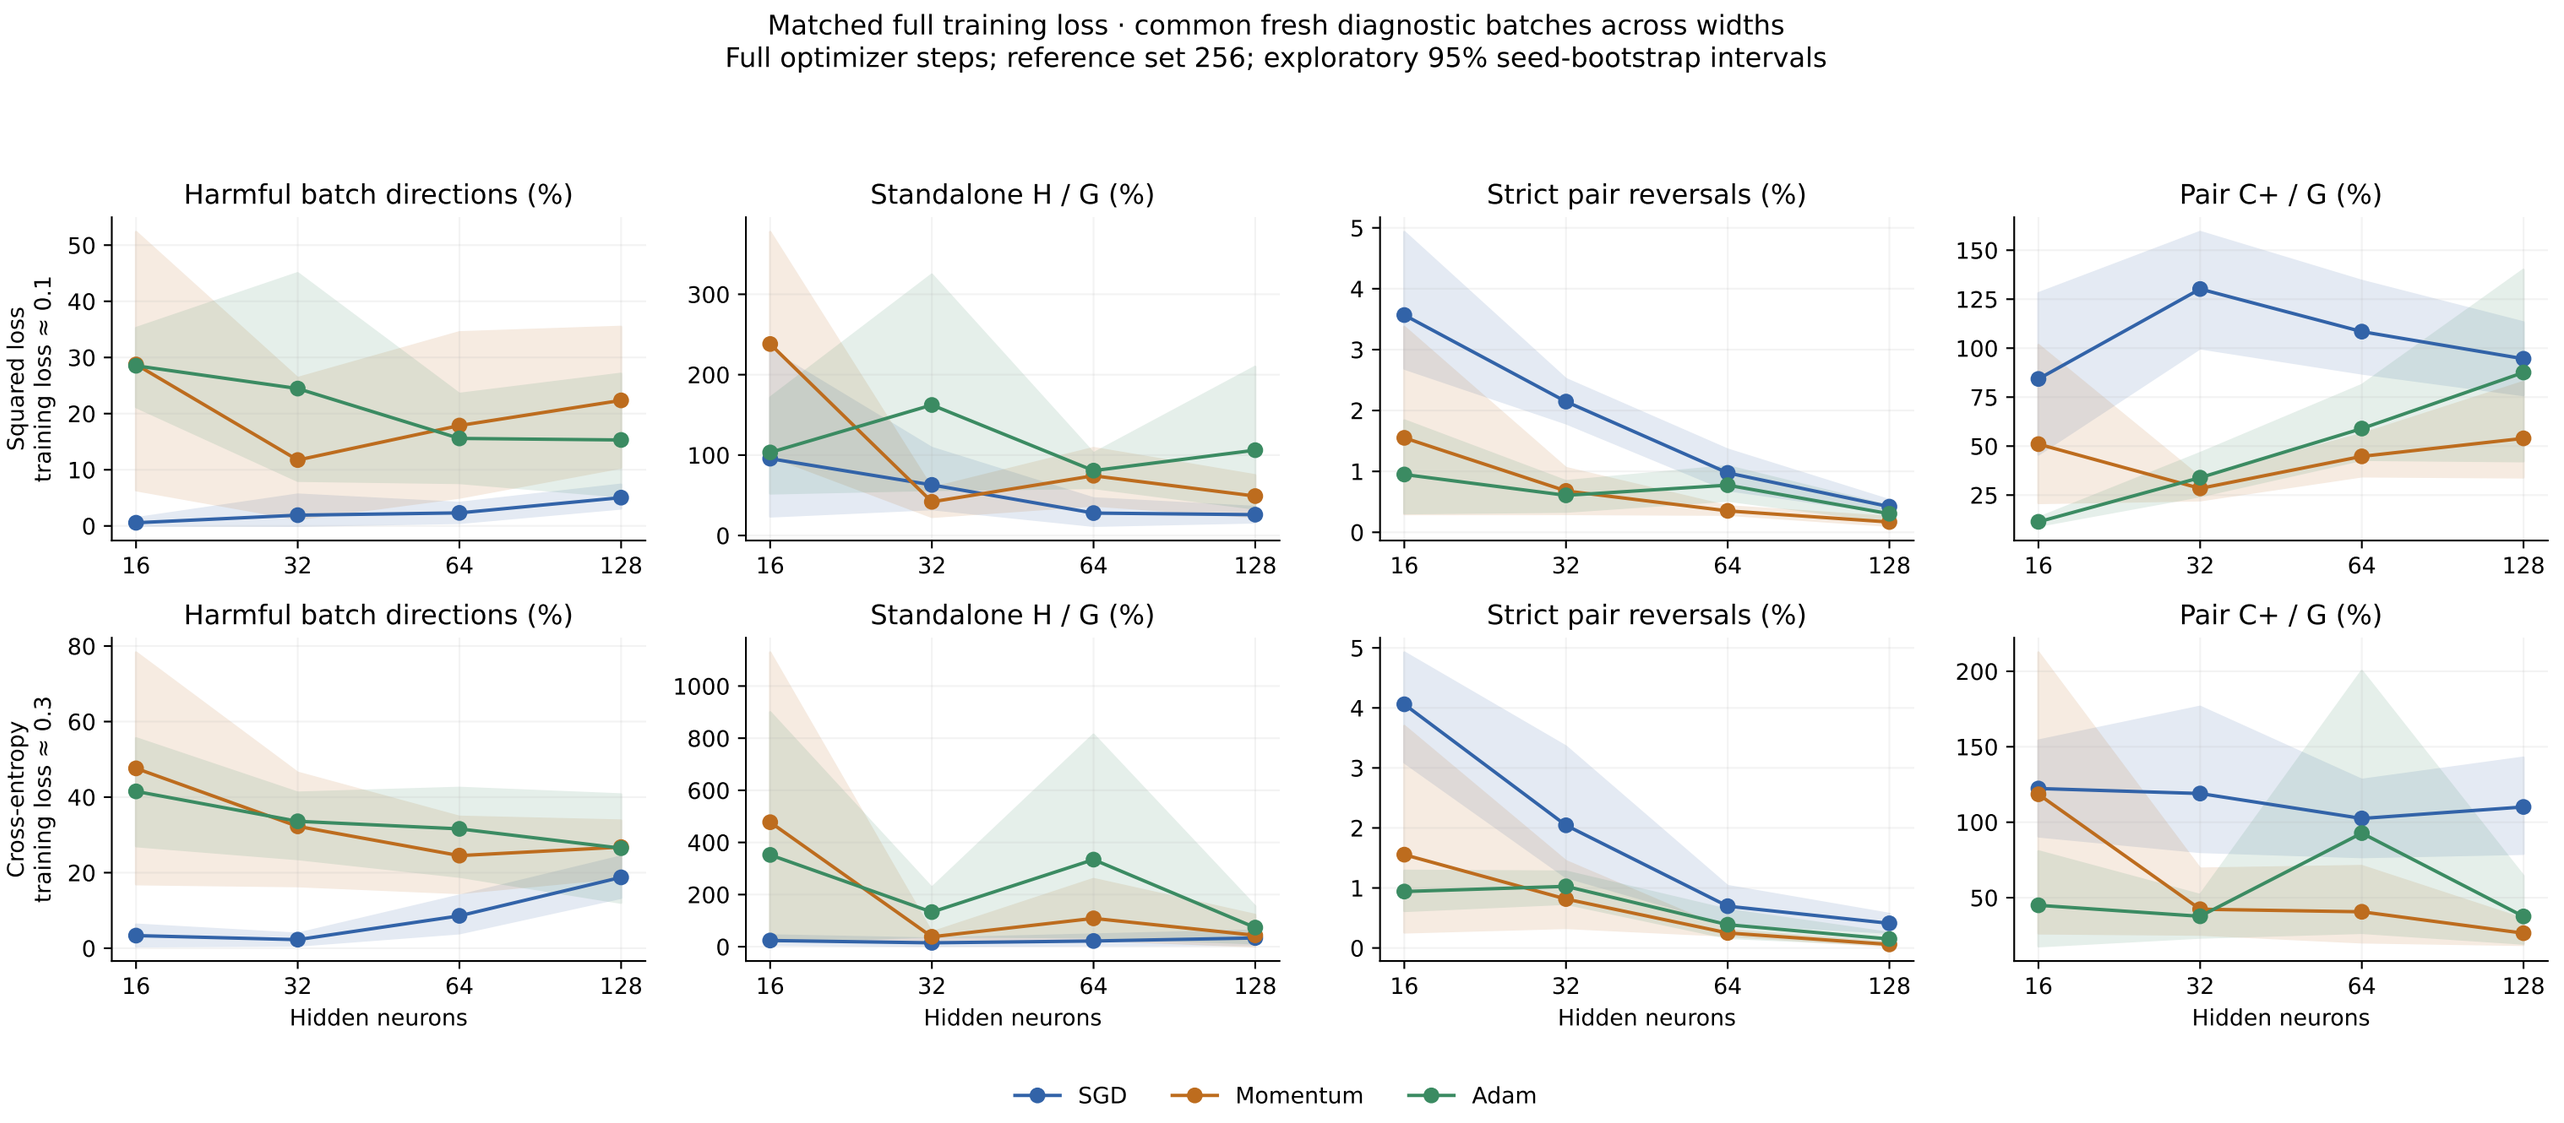
Matched full training loss · common fresh diagnostic batches across widths
Full optimizer steps; reference set 256; exploratory 95% seed-bootstrap intervals
Harmful batch directions (%)
Standalone H / G (%)
Strict pair reversals (%)
Pair C+ / G (%)
Squared loss
training loss ≈ 0.1
Cross-entropy
training loss ≈ 0.3
16
32
64
128
Hidden neurons
SGD
Momentum
Adam