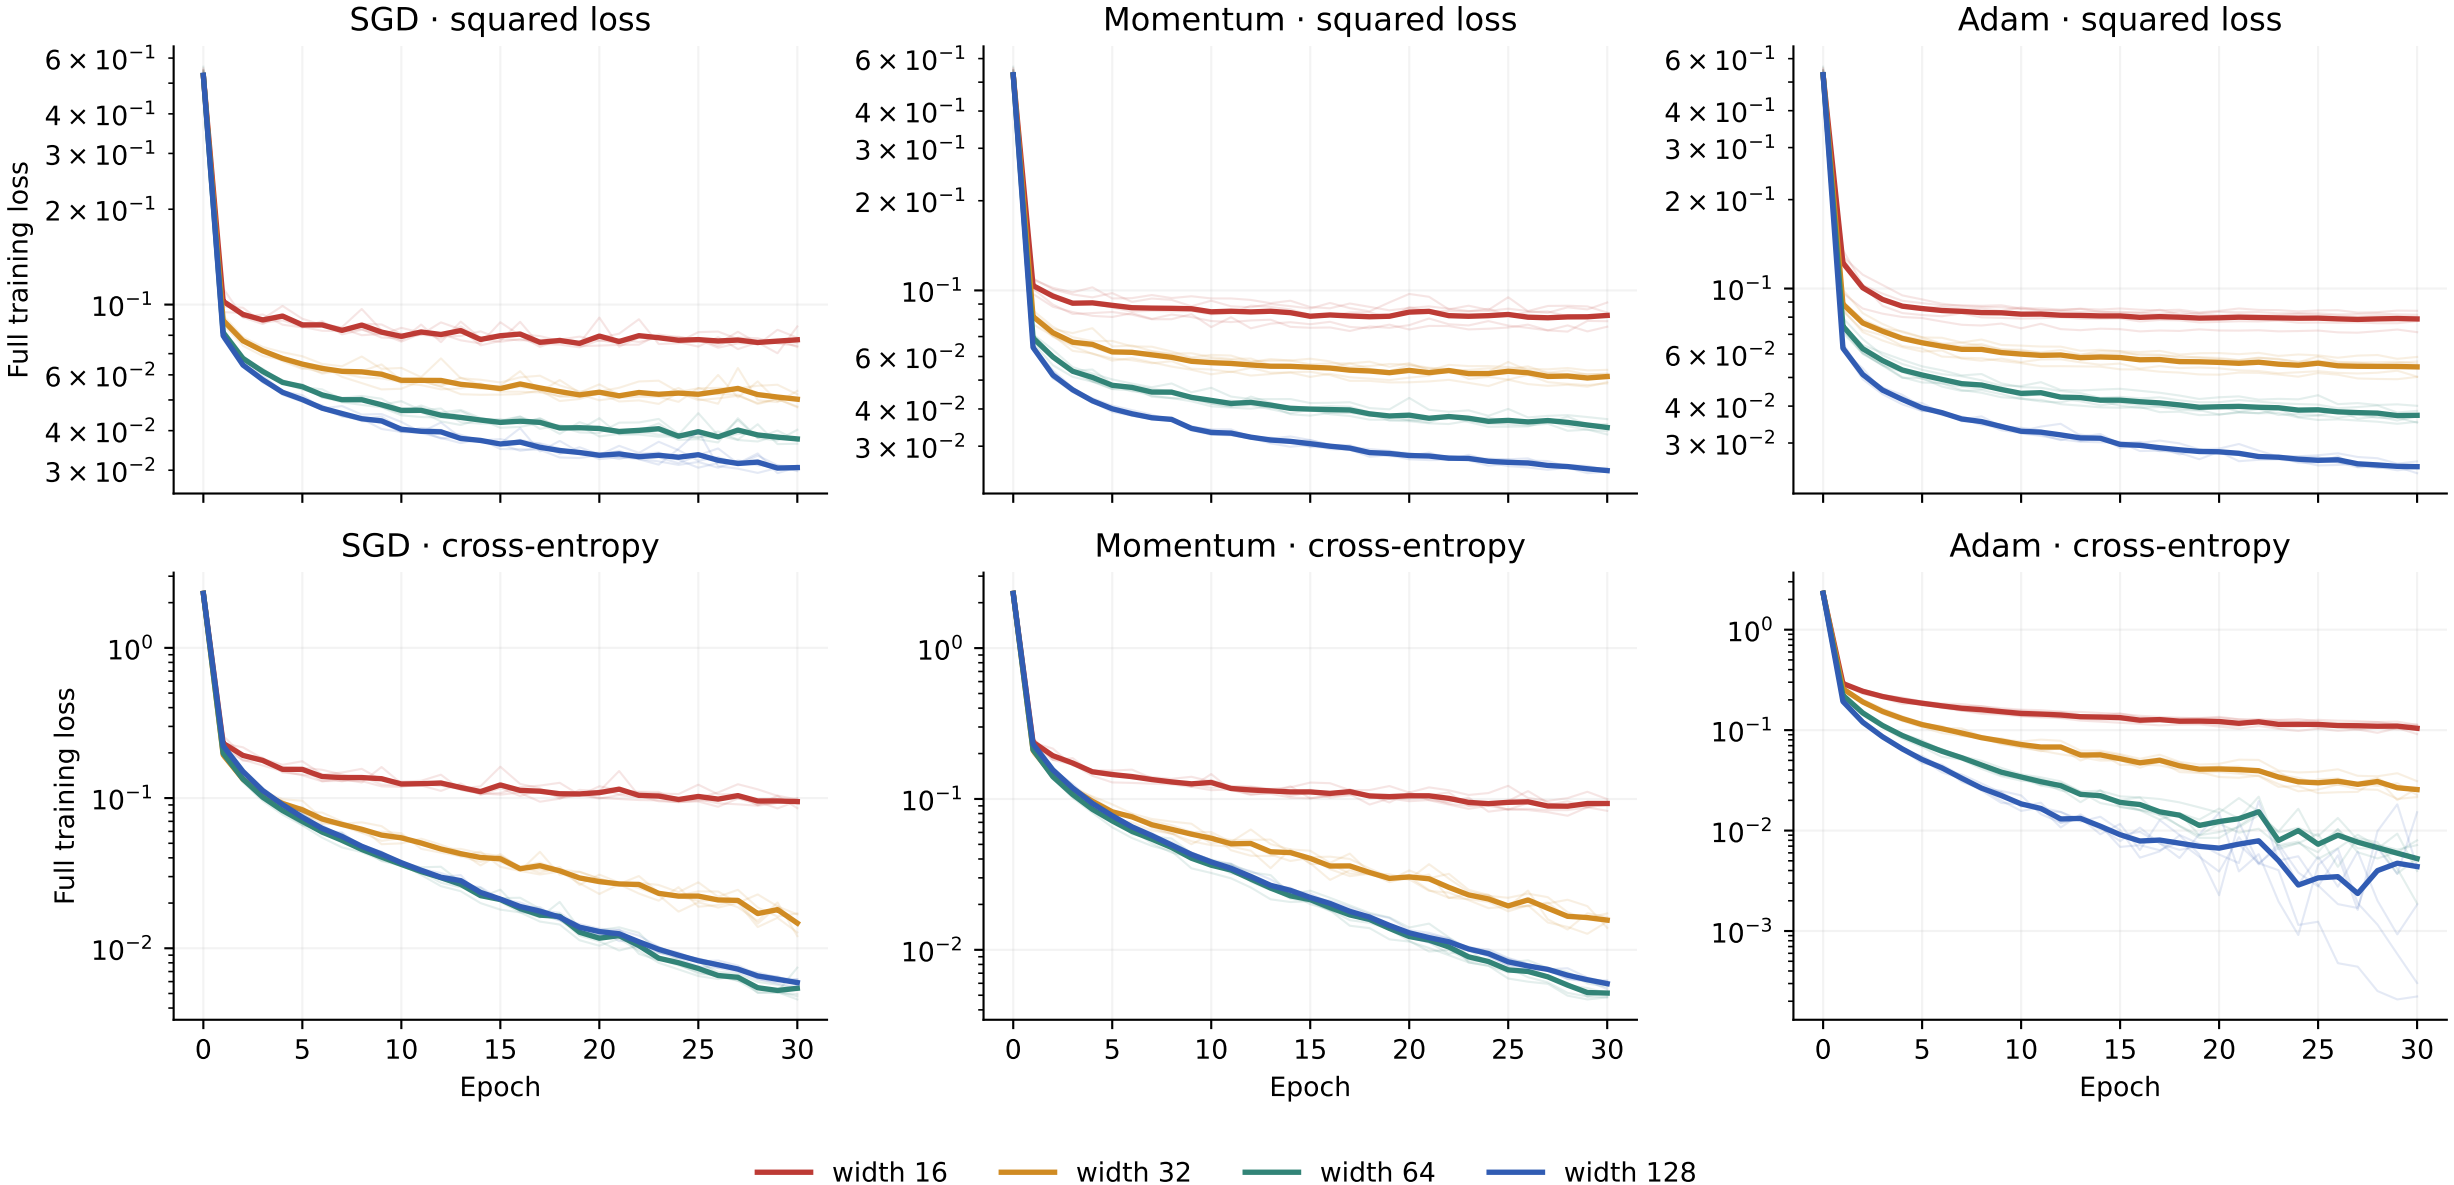
MNIST: one hidden ReLU layer · five paired seeds · constant learning rates
SGD · squared loss
Momentum · squared loss
Adam · squared loss
SGD · cross-entropy
Momentum · cross-entropy
Adam · cross-entropy
Full training loss
Epoch
width 16
width 32
width 64
width 128